

## 4. Pretained GloVe with Bi-LSTM

November 4, 2020

Acknowledgements - <https://www.kaggle.com/lakshmi25npathi/imdb-dataset-of-50k-movie-reviews/notebooks>

```
[ ]: from google.colab import drive
drive.mount('/content/gdrive')
```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force\_remount=True).

```
[ ]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O
import seaborn as sns # data visualization
import matplotlib.pyplot as plt # data visualization
from nltk.corpus import stopwords # text preprocessing - stopwords
import re, string, unicodedata # text preprocessing - regular expressions,
    ↪ string
from bs4 import BeautifulSoup # html processing
from wordcloud import WordCloud, STOPWORDS # visualizing word cloud from corpus,
    ↪ & ignoring stopwords
from collections import Counter # counter for most common words
from sklearn.feature_extraction.text import CountVectorizer # feature-oriented,
    ↪ counting of words
from sklearn.model_selection import train_test_split # splitting the dataset,
    ↪ into train & test sets
from keras.preprocessing import text, sequence # word tokenization
from keras.models import Sequential # class to construct the model
from keras.layers import Embedding, Bidirectional, LSTM, Dense, Dropout,
    ↪ Flatten # RNN layers to be used
from keras.optimizers import Adam # optimizer to be used
from keras.callbacks import ReduceLROnPlateau # learning rate decay on plateau
```

```
[ ]: !pip install -q kaggle
```

```
[ ]: from google.colab import files
```

```
[ ]: files.upload()
```

<IPython.core.display.HTML object>

Saving kaggle.json to kaggle (1).json

```
[ ]: {'kaggle.json':  
      b'{"username":"nabhsanjaymehtautd","key":"ee0e2e2e8b50d345f23e44404b090088"}'}
```

```
[ ]: !mkdir ~/.kaggle/
```

mkdir: cannot create directory '/root/.kaggle/': File exists

```
[ ]: !cp kaggle.json ~/.kaggle/
```

```
[ ]: !chmod 60 ~/.kaggle/kaggle.json
```

```
[ ]: !kaggle datasets list
```

Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600 /root/.kaggle/kaggle.json'

Warning: Looks like you're using an outdated API Version, please consider updating (server 1.5.9 / client 1.5.4)

ref	size	lastUpdated	downloadCount	title
terenceshin/covid19s-impact-on-airport-traffic				COVID-19's Impact on Airport Traffic
	106KB	2020-10-19 12:40:17	1885	
sootersaalu/amazon-top-50-bestselling-books-2009-2019				Amazon Top 50 Bestselling Books 2009 - 2019
	15KB	2020-10-13 09:39:21	1810	
thomaskonstantin/highly-rated-children-books-and-stories				Highly Rated Children Books And Stories
	106KB	2020-10-24 12:09:59	379	
tunguz/euro-parliament-proceedings-1996-2011				Euro Parliament Proceedings 1996 - 2011
	1GB	2020-10-26 17:48:29	26	
rishidamarla/judicial-expenditures-across-all-50-states				Judicial Expenditures across all 50 States
	2KB	2020-10-25 00:07:45	235	
docstein/brics-world-bank-indicators				BRICS World Bank Indicators
	4MB	2020-10-22 12:18:40	397	
kanishk307/6000-indian-food-recipes-dataset				6000+ Indian Food Recipes Dataset
	9MB	2020-10-24 01:08:23	475	
elvinagammed/chatbots-intent-recognition-dataset				Chatbots: Intent Recognition Dataset
	17KB	2020-10-23 07:44:59	196	
omarhanny/500-greatest-songs-of-all-time				500 Greatest Songs of All Time
	33KB	2020-10-26 13:36:09	484	
balraj98/synthetic-objective-testing-set-sots-reside				Synthetic Objective Testing Set (SOTS) [RESIDE]
	415MB	2020-10-24 10:07:29	53	
lunamcbride24/pokemon-type-matchup-data				Pokemon Type Matchup Data
	9KB	2020-10-14 18:56:23	288	

gaurav2796/kaggle-competitions-rankings-and-kernels	Kaggle Competitions, Rankings and Kernels	698KB	2020-10-15 04:05:15	37
balraj98/indoor-training-set-its-residestandard	Indoor Training Set (ITS) [RESIDE-Standard]	5GB	2020-10-24 10:07:30	37
romazepa/moscow-schools-winners-of-educational-olympiads	Moscow schools - winners of educational Olympiads	1MB	2020-10-12 21:45:01	58
sootersaalu/nigerian-songs-spotify	Nigerian Songs Spotify	24KB	2020-10-25 19:10:23	68
salmaneunus/mechanical-tools-dataset	Mechanical Tools Classification Dataset	652MB	2020-11-01 11:28:22	178
thanatoz/hinglish-blogs	Hinglish blogs	2MB	2020-10-13 18:16:05	12
shivamb/netflix-shows	Netflix Movies and TV Shows	971KB	2020-01-20 07:33:56	53154
nehaprabhavalkar/indian-food-101	Indian Food 101	7KB	2020-09-30 06:23:43	6057
heeraldedhia/groceries-dataset	Groceries dataset	257KB	2020-09-17 04:36:08	6282

```
[ ]: !kaggle competitions download -c transferlearning-dl-spring2020
```

```
Warning: Your Kaggle API key is readable by other users on this system! To fix
this, you can run 'chmod 600 /root/.kaggle/kaggle.json'
Warning: Looks like you're using an outdated API Version, please consider
updating (server 1.5.9 / client 1.5.4)
train.csv.zip: Skipping, found more recently modified local copy (use --force to
force download)
test.csv: Skipping, found more recently modified local copy (use --force to
force download)
sample_submission.csv: Skipping, found more recently modified local copy (use
--force to force download)
```

```
[ ]: !unzip train.csv.zip -d train
```

```
Archive: train.csv.zip
replace train/train.csv? [y]es, [n]o, [A]ll, [N]one, [r]ename: y
inflating: train/train.csv
```

```
[ ]: data = pd.read_csv('/content/train/train.csv', encoding = "ISO-8859-1")
testdata = pd.read_csv('/content/test.csv', encoding="ISO-8859-1")
```

```
[ ]: data.head()
```


```
[ ]:
      id      text  target
0  86426  @USER She should ask a few native Americans wh...      1
1  16820  Amazon is investigating Chinese employees who ...      0
2  62688  @USER Someone should'veTaken" this piece of sh...      1
3  43605  @USER @USER Obama wanted liberals & illeg...
```

```
[ ]: sns.set_style('darkgrid')
sns.countplot(data.target)
```

```
/usr/local/lib/python3.6/dist-packages/seaborn/_decorators.py:43: FutureWarning:
Pass the following variable as a keyword arg: x. From version 0.12, the only
valid positional argument will be `data`, and passing other arguments without an
explicit keyword will result in an error or misinterpretation.
```

FutureWarning

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f8768f4d630>
```



```
[ ]: import nltk
nltk.download('stopwords')
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
```

```
[nltk_data] Package stopwords is already up-to-date!
```

```
[ ]: True
```

```
[ ]: stop = set(stopwords.words('english'))
punctuation = list(string.punctuation)
stop.update(punctuation)
```

```
[ ]: def strip_html(text):
    soup = BeautifulSoup(text, "html.parser")
    return soup.get_text()

    #Removing the square brackets
def remove_between_square_brackets(text):
    return re.sub('\[[^\]]*\]', '', text)
    # Removing URL's
def remove_between_square_brackets(text):
    return re.sub(r'http\S+', '', text)
    #Removing the stopwords from text
def remove_stopwords(text):
    final_text = []
    for i in text.split():
        if i.strip().lower() not in stop and i.strip().lower().isalpha():
            final_text.append(i.strip().lower())
    return " ".join(final_text)
    #Removing the noisy text
def denoise_text(text):
    text = strip_html(text)
    text = remove_between_square_brackets(text)
    text = remove_stopwords(text)
    return text
    #Apply function on review column
data['text']=data['text'].apply(denoise_text)
```

```
[ ]: data.head()
```

```
[ ]:      id                                text  target
0   86426                        ask native americans take      1
1   16820  amazon investigating chinese employees selling...      0
2   62688                        someone piece shit      1
3   43605  obama wanted liberals illegals move red states      0
4   97670                        liberals kookoo      1
```

```
[ ]: fig, (ax1, ax2) = plt.subplots(1,2, figsize=(12, 8))
text_len = data[data['target'] == 1]['text'].str.len()
ax1.set_title('Positive Reviews')
ax1.hist(text_len, color='green')
text_len = data[data['target'] == 0]['text'].str.len()
ax2.set_title('Negative Reviews')
ax2.hist(text_len, color='red')
fig.suptitle('Character Count in Reviews')
```

```
[ ]: Text(0.5, 0.98, 'Character Count in Reviews')
```

Character Count in Reviews




```
[ ]: def get_top_ngrams(corpus, n, g):

    vec = CountVectorizer(ngram_range=(g, g)).fit(corpus)
    bag_of_words = vec.transform(corpus)
    sum_words = bag_of_words.sum(axis=0)
    words_freq = [(word, sum_words[0, idx]) for word, idx in vec.vocabulary_.
    ↪items()]
    words_freq = sorted(words_freq, key = lambda x: x[1], reverse=True)
    return words_freq[:n]
```

```
[ ]: plt.figure(figsize=(16, 9))
most_common_uni = dict(get_top_ngrams(data.text, 20, 3))
sns.barplot(x=list(most_common_uni.values()), y=list(most_common_uni.keys()))
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f8737dfdef0>
```



```
[ ]: X_train, X_test, y_train, y_test = train_test_split(data.text, data.target,
                                                    train_size=0.9,
                                                    random_state=0)

X_test_temp = X_test
y_test_temp = y_test
```

```
[ ]: max_features = 10000
max_len = 128

tokenizer = text.Tokenizer(num_words=max_features)
tokenizer.fit_on_texts(X_train)

tokenized_train = tokenizer.texts_to_sequences(X_train)
X_train = sequence.pad_sequences(tokenized_train, maxlen=max_len)

tokenized_test = tokenizer.texts_to_sequences(X_test)
X_test = sequence.pad_sequences(tokenized_test, maxlen=max_len)
```

```
[ ]: !wget http://nlp.stanford.edu/data/glove.6B.zip
```

```
--2020-11-03 20:37:41-- http://nlp.stanford.edu/data/glove.6B.zip
Resolving nlp.stanford.edu (nlp.stanford.edu)... 171.64.67.140
Connecting to nlp.stanford.edu (nlp.stanford.edu)|171.64.67.140|:80...
connected.
HTTP request sent, awaiting response... 302 Found
Location: https://nlp.stanford.edu/data/glove.6B.zip [following]
--2020-11-03 20:37:41-- https://nlp.stanford.edu/data/glove.6B.zip
Connecting to nlp.stanford.edu (nlp.stanford.edu)|171.64.67.140|:443...
```

```

connected.
HTTP request sent, awaiting response... 301 Moved Permanently
Location: http://downloads.cs.stanford.edu/nlp/data/glove.6B.zip [following]
--2020-11-03 20:37:41-- http://downloads.cs.stanford.edu/nlp/data/glove.6B.zip
Resolving downloads.cs.stanford.edu (downloads.cs.stanford.edu)... 171.64.64.22
Connecting to downloads.cs.stanford.edu
(downloads.cs.stanford.edu)|171.64.64.22|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 862182613 (822M) [application/zip]
Saving to: 'glove.6B.zip.2'

```

```

glove.6B.zip.2      100%[=====>] 822.24M  2.13MB/s    in 6m 32s

```

```

2020-11-03 20:44:13 (2.10 MB/s) - 'glove.6B.zip.2' saved [862182613/862182613]

```

```
[ ]: !unzip glove*.zip
```

```

Archive:  glove.6B.zip
replace glove.6B.50d.txt? [y]es, [n]o, [A]ll, [N]one, [r]ename: y
  inflating: glove.6B.50d.txt
replace glove.6B.100d.txt? [y]es, [n]o, [A]ll, [N]one, [r]ename: y
  inflating: glove.6B.100d.txt
replace glove.6B.200d.txt? [y]es, [n]o, [A]ll, [N]one, [r]ename: y
  inflating: glove.6B.200d.txt
replace glove.6B.300d.txt? [y]es, [n]o, [A]ll, [N]one, [r]ename: y
  inflating: glove.6B.300d.txt

```

```
[ ]: !ls
      !pwd
```

```

gdrive          glove.6B.50d.txt  'kaggle (1).json'      test.csv
glove.6B.100d.txt glove.6B.zip      kaggle.json            train
glove.6B.200d.txt glove.6B.zip.1     sample_data            train.csv.zip
glove.6B.300d.txt glove.6B.zip.2     sample_submission.csv
/content

```

```
[ ]: EMBEDDING_FILE = "/content/glove.6B.300d.txt"
```

```
[ ]: def get_coeffs(word, *arr):
      return word, np.asarray(arr, dtype='float32')
      embeddings_dict = dict(get_coeffs(*o.rstrip().rsplit(' ')) for o in
      ↪open(EMBEDDING_FILE))

```

```
[ ]: all_embs = np.stack(embeddings_dict.values())
      emb_mean, emb_std = all_embs.mean(), all_embs.std()
      embed_size = all_embs.shape[1]

```



```

word_index = tokenizer.word_index
num_words = min(max_features, len(word_index))

embedding_matrix = np.random.normal(emb_mean, emb_std, (num_words, embed_size))

for word, i in word_index.items():

    if i >= num_words: continue
    embedding_vector = embeddings_dict.get(word)
    if embedding_vector is not None: embedding_matrix[i] = embedding_vector

```

/usr/local/lib/python3.6/dist-packages/IPython/core/interactiveshell.py:2822:  
FutureWarning: arrays to stack must be passed as a "sequence" type such as list  
or tuple. Support for non-sequence iterables such as generators is deprecated as  
of NumPy 1.16 and will raise an error in the future.

```
if self.run_code(code, result):
```

```
[ ]: batch_size = 256
      epochs=10
      embed_size=300
```

```
[ ]: learning_rate_reduction = ReduceLROnPlateau(monitor='val_loss', patience=2,
↳ verbose=1, factor=0.5, min_lr=0.00001)
```

```
[ ]: from keras import backend as K
```

```
[ ]: def f_score(y_true, y_pred, threshold=0.1, beta=2):

    tp = tp_score(y_true, y_pred, threshold)
    fp = fp_score(y_true, y_pred, threshold)
    fn = fn_score(y_true, y_pred, threshold)

    precision = tp / (tp + fp)
    recall = tp / (tp + fn)
    return (1+beta**2) * ((precision * recall) / ((beta**2)*precision + recall))

```

```
[ ]: def tp_score(y_true, y_pred, threshold=0.1):

    tp_3d = K.concatenate(
        [
            K.cast(K.expand_dims(K.flatten(y_true)), 'bool'),
            K.cast(K.expand_dims(K.flatten(K.greater(y_pred, K.
↳ constant(threshold))))), 'bool'),
            K.cast(K.ones_like(K.expand_dims(K.flatten(y_pred))), 'bool')
        ], axis=1
    )

```

```

    tp = K.sum(K.cast(K.all(tp_3d, axis=1), 'int32'))

    return tp

def fp_score(y_true, y_pred, threshold=0.1):

    fp_3d = K.concatenate(
        [
            K.cast(K.expand_dims(K.flatten(K.abs(y_true - K.
→ones_like(y_true))))), 'bool'),
            K.cast(K.expand_dims(K.flatten(K.greater(y_pred, K.
→constant(threshold))))), 'bool'),
            K.cast(K.ones_like(K.expand_dims(K.flatten(y_pred))), 'bool')
        ], axis=-1
    )

    fp = K.sum(K.cast(K.all(fp_3d, axis=1), 'int32'))

    return fp

def fn_score(y_true, y_pred, threshold=0.1):

    fn_3d = K.concatenate(
        [
            K.cast(K.expand_dims(K.flatten(y_true))), 'bool'),
            K.cast(K.expand_dims(K.flatten(K.abs(K.cast(K.greater(y_pred, K.
→constant(threshold)), 'float') - K.ones_like(y_pred))))), 'bool'),
            K.cast(K.ones_like(K.expand_dims(K.flatten(y_pred))), 'bool')
        ], axis=1
    )

    fn = K.sum(K.cast(K.all(fn_3d, axis=1), 'int32'))

    return fn

def precision_score(y_true, y_pred, threshold=0.1):

    tp = tp_score(y_true, y_pred, threshold)
    fp = fp_score(y_true, y_pred, threshold)

    return tp / (tp + fp)

def recall_score(y_true, y_pred, threshold=0.1):

```

```

tp = tp_score(y_true, y_pred, threshold)
fn = fn_score(y_true, y_pred, threshold)

return tp / (tp + fn)

```

```

[ ]: model = Sequential()

model.add(Embedding(max_features, output_dim=embed_size,
    ↳weights=[embedding_matrix], input_length=max_len, trainable=False))
model.add(Bidirectional(LSTM(units=128)))
model.add(Dropout(rate=0.8))
model.add(Dense(units=16, activation='relu'))
model.add(Dense(units=1, activation='sigmoid'))

model.compile(optimizer=Adam(lr=0.002), loss='binary_crossentropy',
    metrics=f_score)

```

```

[ ]: model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 128, 300)	3000000
bidirectional (Bidirectional)	(None, 256)	439296
dropout (Dropout)	(None, 256)	0
dense (Dense)	(None, 16)	4112
dense_1 (Dense)	(None, 1)	17
Total params: 3,443,425		
Trainable params: 443,425		
Non-trainable params: 3,000,000		

```

[ ]: history = model.fit(X_train, y_train, epochs=epochs, batch_size=batch_size,
    ↳validation_data=(X_test, y_test), callbacks=[learning_rate_reduction])

```

Epoch 1/10

33/33 [=====] - 2s 73ms/step - loss: 0.6321 - f\_score: 0.7129 - val\_loss: 0.5864 - val\_f\_score: 0.7172

Epoch 2/10

33/33 [=====] - 2s 48ms/step - loss: 0.5678 - f\_score:

```

0.7190 - val_loss: 0.5355 - val_f_score: 0.7180
Epoch 3/10
33/33 [=====] - 2s 48ms/step - loss: 0.5366 - f_score:
0.7190 - val_loss: 0.5132 - val_f_score: 0.7203
Epoch 4/10
33/33 [=====] - 2s 48ms/step - loss: 0.5190 - f_score:
0.7280 - val_loss: 0.5102 - val_f_score: 0.7290
Epoch 5/10
33/33 [=====] - 2s 48ms/step - loss: 0.5085 - f_score:
0.7302 - val_loss: 0.5158 - val_f_score: 0.7342
Epoch 6/10
33/33 [=====] - 2s 48ms/step - loss: 0.4942 - f_score:
0.7346 - val_loss: 0.5067 - val_f_score: 0.7297
Epoch 7/10
33/33 [=====] - 2s 48ms/step - loss: 0.4860 - f_score:
0.7389 - val_loss: 0.5085 - val_f_score: 0.7313
Epoch 8/10
33/33 [=====] - ETA: 0s - loss: 0.4653 - f_score:
0.7451
Epoch 00008: ReduceLROnPlateau reducing learning rate to 0.0010000000474974513.
33/33 [=====] - 2s 48ms/step - loss: 0.4653 - f_score:
0.7451 - val_loss: 0.5075 - val_f_score: 0.7339
Epoch 9/10
33/33 [=====] - 2s 48ms/step - loss: 0.4461 - f_score:
0.7519 - val_loss: 0.5250 - val_f_score: 0.7306
Epoch 10/10
33/33 [=====] - ETA: 0s - loss: 0.4260 - f_score:
0.7619
Epoch 00010: ReduceLROnPlateau reducing learning rate to 0.0005000000237487257.
33/33 [=====] - 2s 48ms/step - loss: 0.4260 - f_score:
0.7619 - val_loss: 0.5382 - val_f_score: 0.7306

```

```

[ ]: print("Model Accuracy on Training Data: ", round(model.evaluate(X_train,
    ↪y_train)[1]*100), "%")
print("Model Accuracy on Testing Data: ", round(model.evaluate(X_test,
    ↪y_test)[1]*100), "%")

```

```

263/263 [=====] - 2s 8ms/step - loss: 0.3937 - f_score:
0.7533
Model Accuracy on Training Data: 75 %
30/30 [=====] - 0s 7ms/step - loss: 0.5382 - f_score:
0.7171
Model Accuracy on Testing Data: 72 %

```

```

[ ]: history.history

```

```
[ ]: {'f_score': [0.7128528952598572,
0.7189640998840332,
0.7189900875091553,
0.7280216217041016,
0.7301872372627258,
0.7346366047859192,
0.7388955354690552,
0.745067298412323,
0.7518593072891235,
0.7619423270225525],
'loss': [0.632097601890564,
0.5678114891052246,
0.5365914106369019,
0.5190249681472778,
0.5085451006889343,
0.49418380856513977,
0.48599594831466675,
0.46534425020217896,
0.4460606575012207,
0.4260425567626953],
'lr': [0.002, 0.002, 0.002, 0.002, 0.002, 0.002, 0.002, 0.002, 0.001, 0.001],
'val_f_score': [0.7171559929847717,
0.7179653644561768,
0.7203402519226074,
0.729023814201355,
0.7342402338981628,
0.7296569347381592,
0.7313453555107117,
0.7338710427284241,
0.7306410074234009,
0.7306411862373352],
'val_loss': [0.5863938331604004,
0.5354620218276978,
0.5131537318229675,
0.5101721882820129,
0.515751302242279,
0.5066903233528137,
0.5085200667381287,
0.5074759721755981,
0.5249583721160889,
0.5382400751113892]}
```


```
[ ]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(18, 8))
ax1.plot(history.history['f_score'])
ax1.plot(history.history['val_f_score'])
ax1.set_title('Training & Testing Accuracy')
ax1.set_xlabel('Epochs')
```

```

ax1.set_ylabel('Accuracy')
ax1.legend(['accuracy', 'val_f_score'])
ax2.plot(history.history['loss'])
ax2.plot(history.history['val_loss'])
ax2.set_title('Training & Testing Loss')
ax1.set_xlabel('Epochs')
ax1.set_ylabel('Loss')
ax2.legend(['loss', 'val_loss'])

```

```
[ ]: <matplotlib.legend.Legend at 0x7f86cec5df98>
```



```
[ ]: testdata['text']=testdata['text'].apply(denoise_text)
```

```
[ ]: testdata_X = testdata['text']
```

```
[ ]: testdata_X.head()
```

```

[ ]: 0          go home url
      1          oh tough
      2  canada need another already enough great
      3          scare every playing hockey warped
      4          throwing bullshit flag url
      Name: text, dtype: object

```

```
[ ]:
```

```

[ ]: max_features = 10000
      max_len = 128

```

```
tokenizer = text.Tokenizer(num_words=max_features)
tokenizer.fit_on_texts(testdata_X)

tokenized_train = tokenizer.texts_to_sequences(testdata_X)
testdata_X = sequence.pad_sequences(tokenized_train, maxlen=max_len)
```

```
[ ]: predictions = model.predict(testdata_X)
```

```
[ ]: type(predictions)
```

```
[ ]: numpy.ndarray
```

```
[ ]: submit = pd.DataFrame(data=predictions, columns=['target'])
```

```
[ ]: submit
```

```
[ ]:
      target
0      0.076413
1      0.088850
2      0.309796
3      0.425224
4      0.056183
...
3889  0.150014
3890  0.072804
3891  0.651642
3892  0.563634
3893  0.096666
```

```
[3894 rows x 1 columns]
```

```
[ ]: submit1 = pd.DataFrame(data=testdata['id'], columns=['id'])
```

```
[ ]: submit1
```

```
[ ]:
      id
0     90194
1     77444
2     13384
3     54920
4     56117
...
3889  90041
3890  98824
3891  95338
3892  67210
3893  46552
```

[3894 rows x 1 columns]

```
[ ]: submit1['target'] = submit['target']
```

```
[ ]: submit1.to_csv('/content/gdrive/My Drive/submit2.csv', index = False)
```

```
[ ]: submit1
```

```
[ ]:
      id  target
0   90194  0.076413
1   77444  0.088850
2   13384  0.309796
3   54920  0.425224
4   56117  0.056183
...    ...    ...
3889  90041  0.150014
3890  98824  0.072804
3891  95338  0.651642
3892  67210  0.563634
3893  46552  0.096666
```

[3894 rows x 2 columns]

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```



[ ]:	
[ ]:	
[ ]:	
[ ]:	
[ ]:	
[ ]:	